

Import necessary dependencies

```
In [2]: import pandas as pd
import numpy as np
# import text_normalizer as tn
import model_evaluation_utils as meu
import importlib

np.set_printoptions(precision=2, linewidth=80)
```

Load and normalize data

```
In [3]: dataset = pd.read_csv('../data/movie_reviews_cleaned.csv')

# take a peek at the data
print(dataset.head())
reviews = np.array(dataset['review'])
sentiments = np.array(dataset['sentiment'])

# build train and test datasets
train_reviews = reviews[:5000]
train_sentiments = sentiments[:5000]
test_reviews = reviews[5000:7000]
test_sentiments = sentiments[5000:7000]

# normalize datasets
# norm_train_reviews = tn.normalize_corpus(train_reviews)
# norm_test_reviews = tn.normalize_corpus(test_reviews)

norm_train_reviews = train_reviews
norm_test_reviews = test_reviews
```

	review	sentiment
0	not bother think would see movie great supspen...	negative
1	careful one get mitt change way look kung fu f...	positive
2	chili palmer tired movie know want success mus...	negative
3	follow little know 1998 british film make budg...	positive
4	dark angel cross huxley brave new world percys...	positive

Traditional Supervised Machine Learning Models

Feature Engineering

```
In [4]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
```

```
# build BOW features on train reviews
cv = CountVectorizer(binary=False, min_df=0.0, max_df=1.0, ngram_range=(1,2))
cv_train_features = cv.fit_transform(norm_train_reviews)
# build TFIDF features on train reviews
tv = TfidfVectorizer(use_idf=True, min_df=0.0, max_df=1.0, ngram_range=(1,2),
                    sublinear_tf=True)
tv_train_features = tv.fit_transform(norm_train_reviews)
```

```
In [5]: # transform test reviews into features
cv_test_features = cv.transform(norm_test_reviews)
tv_test_features = tv.transform(norm_test_reviews)
```

```
In [6]: print('BOW model:> Train features shape:', tv_train_features.shape, ' Test f
print('TFIDF model:> Train features shape:', tv_train_features.shape, ' Test
```

```
BOW model:> Train features shape: (5000, 434563) Test features shape: (200
0, 434563)
TFIDF model:> Train features shape: (5000, 434563) Test features shape: (2
000, 434563)
```

Model Training, Prediction and Performance Evaluation

```
In [7]: from sklearn.linear_model import SGDClassifier, LogisticRegression
```

```
lr = LogisticRegression(penalty='l2', max_iter=100, C=1)
svm = SGDClassifier(loss='hinge', max_iter=100)
```

```
In [8]: # Logistic Regression model on BOW features
```

```
importlib.reload(meu)

lr_bow_predictions = meu.train_predict_model(classifier=lr,
                                             train_features=cv_train_features,
                                             test_features=cv_test_features,
meu.display_model_performance_metrics(true_labels=test_sentiments, predicted
                                     classes=['positive', 'negative'])
```

```
/Users/canxiuzhang/anaconda3/envs/nlp_3_9/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
```

Model Performance metrics:

Accuracy: 0.8605
 Precision: 0.8606
 Recall: 0.8605
 F1 Score: 0.8605

Model Classification report:

	precision	recall	f1-score	support
positive	0.85	0.86	0.86	981
negative	0.87	0.86	0.86	1019
accuracy			0.86	2000
macro avg	0.86	0.86	0.86	2000
weighted avg	0.86	0.86	0.86	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	846	135
negative	144	875

```
In [9]: # Logistic Regression model on TF-IDF features
lr_tfidf_predictions = meu.train_predict_model(classifier=lr,
                                                train_features=tv_train_features,
                                                test_features=tv_test_features,
                                                meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_labels=lr_tfidf_predictions,
                                                classes=['positive', 'negative']))
```

Model Performance metrics:

Accuracy: 0.866
 Precision: 0.8661
 Recall: 0.866
 F1 Score: 0.866

Model Classification report:

	precision	recall	f1-score	support
positive	0.87	0.85	0.86	981
negative	0.86	0.88	0.87	1019
accuracy			0.87	2000
macro avg	0.87	0.87	0.87	2000
weighted avg	0.87	0.87	0.87	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	838	143
negative	125	894

```
In [10]: svm_bow_predictions = meu.train_predict_model(classifier=svm,
                                                         train_features=cv_train_features,
                                                         test_features=cv_test_features,
                                                         meu.display_model_performance_metrics(true_labels=test_sentiments, predicted=svm_bow_predictions,
                                                         classes=['positive', 'negative']))
```

Model Performance metrics:

Accuracy: 0.8405
 Precision: 0.8405
 Recall: 0.8405
 F1 Score: 0.8405

Model Classification report:

	precision	recall	f1-score	support
positive	0.84	0.84	0.84	981
negative	0.84	0.84	0.84	1019
accuracy			0.84	2000
macro avg	0.84	0.84	0.84	2000
weighted avg	0.84	0.84	0.84	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	820	161
negative	158	861

```
In [11]: svm_tfidf_predictions = meu.train_predict_model(classifier=svm,
                                                         train_features=tv_train_feat
                                                         test_features=tv_test_featur
meu.display_model_performance_metrics(true_labels=test_sentiments, predicted
                                     classes=['positive', 'negative'])
```

Model Performance metrics:

Accuracy: 0.8815
 Precision: 0.8815
 Recall: 0.8815
 F1 Score: 0.8815

Model Classification report:

	precision	recall	f1-score	support
positive	0.88	0.87	0.88	981
negative	0.88	0.89	0.88	1019
accuracy			0.88	2000
macro avg	0.88	0.88	0.88	2000
weighted avg	0.88	0.88	0.88	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	856	125
negative	112	907

Newer Supervised Deep Learning Models

```
In [12]: import gensim
import keras
from keras.models import Sequential
from keras.layers import Dropout, Activation, Dense
from sklearn.preprocessing import LabelEncoder

import spacy
import nltk
from nltk.tokenize.toktok import ToktokTokenizer

tokenizer = ToktokTokenizer()

nlp = spacy.load('en_core_web_sm')
```

Prediction class label encoding

```
In [13]: le = LabelEncoder()
num_classes=2
# tokenize train reviews & encode train labels
tokenized_train = [tokenizer.tokenize(text)
                    for text in norm_train_reviews]
y_tr = le.fit_transform(train_sentiments)
y_train = keras.utils.to_categorical(y_tr, num_classes)
# tokenize test reviews & encode test labels
```

```
tokenized_test = [tokenizer.tokenize(text)
                   for text in norm_test_reviews]
y_ts = le.fit_transform(test_sentiments)
y_test = keras.utils.to_categorical(y_ts, num_classes)
```

```
In [14]: # print class label encoding map and encoded labels
print('Sentiment class label map:', dict(zip(le.classes_, le.transform(le.c
print('Sample test label transformation:\n'+ '-'*35,
      '\nActual Labels:', test_sentiments[:3], '\nEncoded Labels:', y_ts[:3]
      '\nOne hot encoded Labels:\n', y_test[:3]))
```

Sentiment class label map: {'negative': 0, 'positive': 1}

Sample test label transformation:

```
-----
Actual Labels: ['negative' 'negative' 'negative']
Encoded Labels: [0 0 0]
One hot encoded Labels:
[[1. 0.]
 [1. 0.]
 [1. 0.]]
```

Feature Engineering with word embeddings

```
In [15]: # build word2vec model
w2v_num_features = 500
w2v_model = gensim.models.Word2Vec(tokenized_train, vector_size=w2v_num_feat
                                   min_count=10, sample=1e-3)
```

```
In [16]: def averaged_word2vec_vectorizer(corpus, model, num_features):
          vocabulary = set(model.wv.index_to_key)

          def average_word_vectors(words, model, vocabulary, num_features):
              feature_vector = np.zeros((num_features,), dtype="float64")
              nwords = 0.

              for word in words:
                  if word in vocabulary:
                      nwords = nwords + 1.
                      feature_vector = np.add(feature_vector, model.wv[word])
              if nwords:
                  feature_vector = np.divide(feature_vector, nwords)

              return feature_vector

          features = [average_word_vectors(tokenized_sentence, model, vocabulary,
                                           for tokenized_sentence in corpus]
          return np.array(features)
```

```
In [17]: # generate averaged word vector features from word2vec model
avg_wv_train_features = averaged_word2vec_vectorizer(corpus=tokenized_train,
                                                    num_features=500)
avg_wv_test_features = averaged_word2vec_vectorizer(corpus=tokenized_test, num
                                                    num_features=500)

In [18]: # feature engineering with GloVe model
train_nlp = [nlp(item) for item in norm_train_reviews]
train_glove_features = np.array([item.vector for item in train_nlp])

test_nlp = [nlp(item) for item in norm_test_reviews]
test_glove_features = np.array([item.vector for item in test_nlp])

In [19]: print('Word2Vec model:> Train features shape:', avg_wv_train_features.shape,
print('GloVe model:> Train features shape:', train_glove_features.shape, 'T

Word2Vec model:> Train features shape: (5000, 500) Test features shape: (2
000, 500)
GloVe model:> Train features shape: (5000, 96) Test features shape: (2000,
96)
```

Modeling with deep neural networks

Building Deep neural network architecture

```
In [20]: def construct_deepnn_architecture(num_input_features):
dnn_model = Sequential()
dnn_model.add(Dense(512, activation='relu', input_shape=(num_input_featu
dnn_model.add(Dropout(0.2))
dnn_model.add(Dense(512, activation='relu'))
dnn_model.add(Dropout(0.2))
dnn_model.add(Dense(512, activation='relu'))
dnn_model.add(Dropout(0.2))
dnn_model.add(Dense(2))
dnn_model.add(Activation('softmax'))

dnn_model.compile(loss='categorical_crossentropy', optimizer='adam',
metrics=['accuracy'])
return dnn_model
```

```
In [21]: w2v_dnn = construct_deepnn_architecture(num_input_features=500)
```

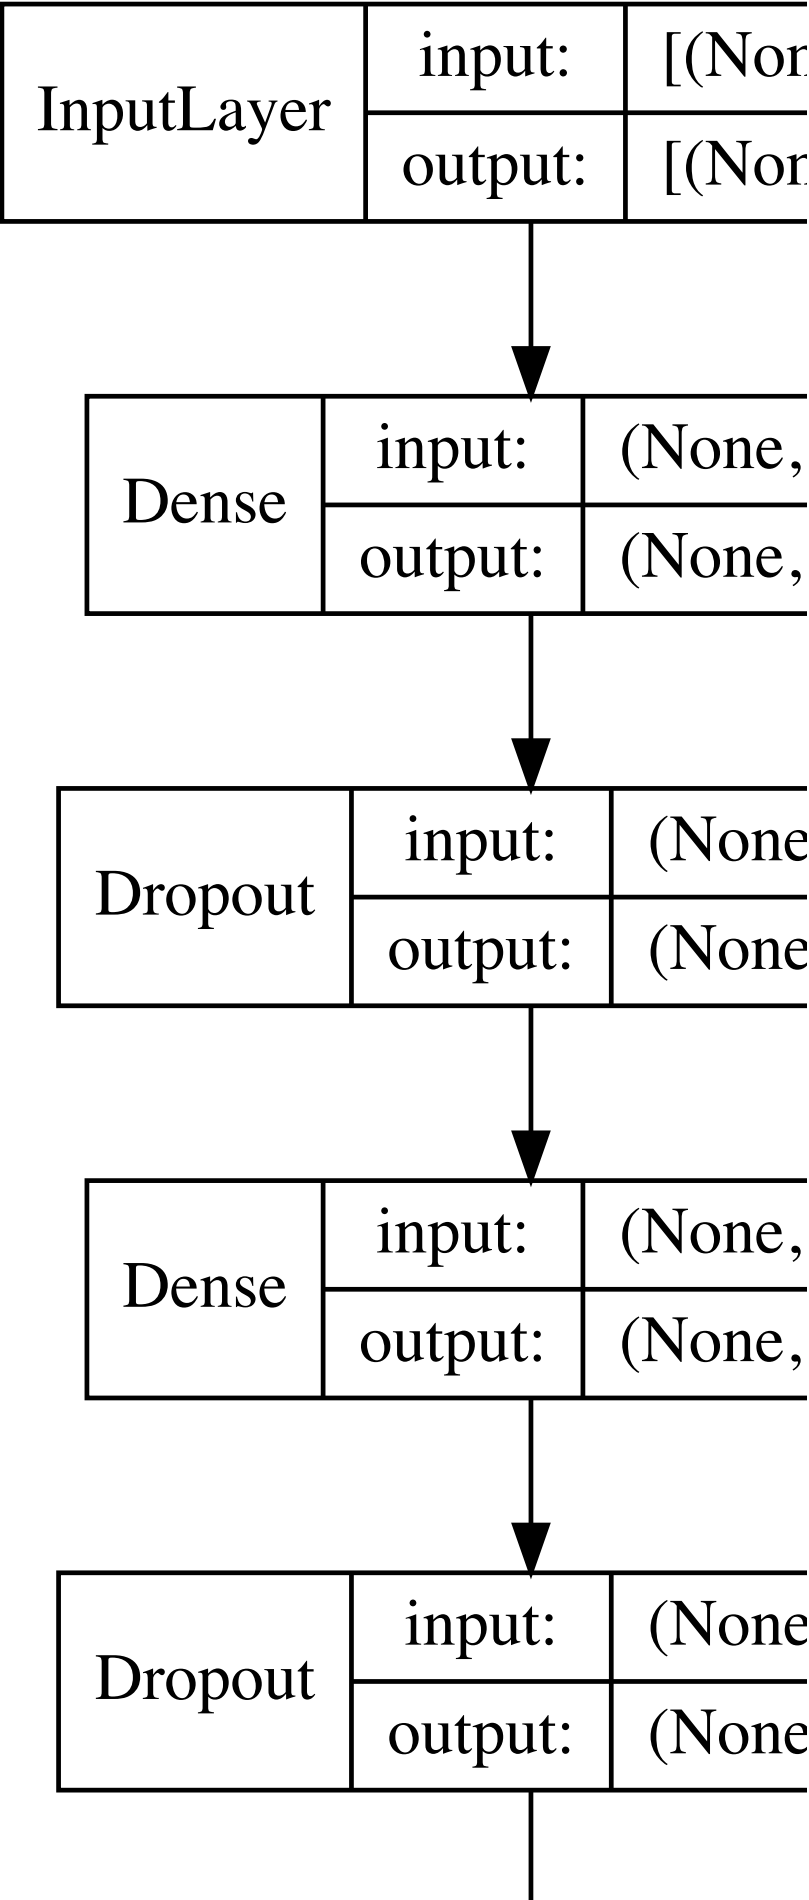
2023-03-18 14:59:10.128574: I tensorflow/core/platform/cpu_feature_guard.c c:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

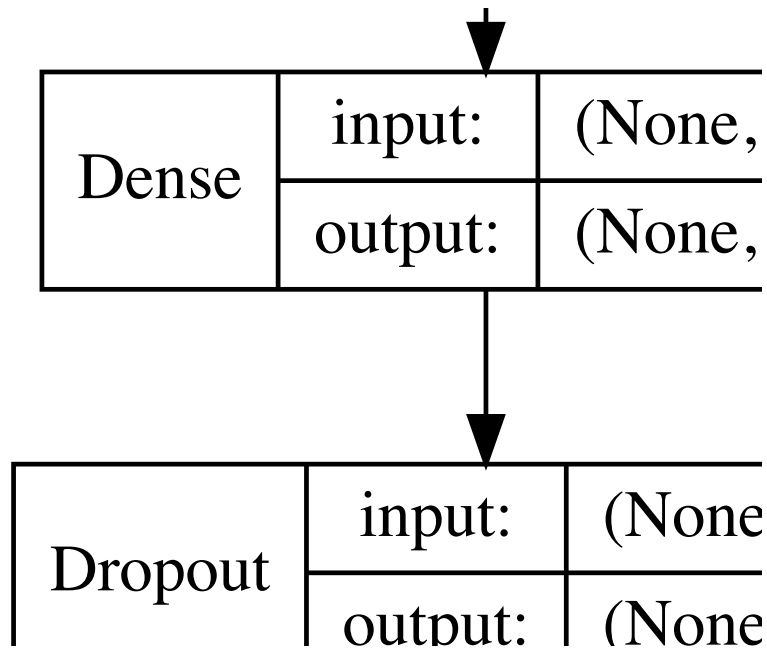
Visualize sample deep architecture


```
In [41]: from IPython.display import SVG
from keras.utils.vis_utils import model_to_dot

SVG(model_to_dot(w2v_dnn, show_shapes=True, show_layer_names=False,
                 rankdir='TB').create(prog='dot', format='svg'))
```

Out [41]:





Model Training, Prediction and Performance Evaluation

```
In [23]: batch_size = 100
w2v_dnn.fit(avg_wv_train_features, y_train, epochs=5, batch_size=batch_size,
            shuffle=True, validation_split=0.1, verbose=1)
```

```
Epoch 1/5
45/45 [=====] - 3s 27ms/step - loss: 0.5057 - accuracy: 0.7553 - val_loss: 0.4687 - val_accuracy: 0.7720
Epoch 2/5
45/45 [=====] - 1s 19ms/step - loss: 0.4599 - accuracy: 0.7816 - val_loss: 0.4587 - val_accuracy: 0.7740
Epoch 3/5
45/45 [=====] - 1s 30ms/step - loss: 0.4508 - accuracy: 0.7862 - val_loss: 0.4510 - val_accuracy: 0.7840
Epoch 4/5
45/45 [=====] - 2s 49ms/step - loss: 0.4475 - accuracy: 0.7869 - val_loss: 0.4551 - val_accuracy: 0.7880
Epoch 5/5
45/45 [=====] - 1s 29ms/step - loss: 0.4386 - accuracy: 0.7956 - val_loss: 0.4412 - val_accuracy: 0.7820
```

```
Out[23]: <keras.callbacks.History at 0x7ff96edc8fa0>
```

```
In [26]: y_pred = w2v_dnn.predict(avg_wv_test_features)
y_classes = np.argmax(y_pred, axis=1)
predictions = le.inverse_transform(y_classes)
```

```
63/63 [=====] - 0s 3ms/step
```

```
In [27]: meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_classes=['positive', 'negative'])
```

Model Performance metrics:

Accuracy: 0.793
 Precision: 0.8005
 Recall: 0.793
 F1 Score: 0.7921

Model Classification report:

	precision	recall	f1-score	support
positive	0.75	0.87	0.80	981
negative	0.85	0.72	0.78	1019
accuracy			0.79	2000
macro avg	0.80	0.79	0.79	2000
weighted avg	0.80	0.79	0.79	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	850	131
negative	283	736

```
In [31]: glove_dnn = construct_deepnn_architecture(num_input_features=96)
```

```
In [32]: batch_size = 100
glove_dnn.fit(train_glove_features, y_train, epochs=5, batch_size=batch_size,
              shuffle=True, validation_split=0.1, verbose=1)
```

Epoch 1/5

45/45 [=====] - 3s 27ms/step - loss: 0.6729 - accuracy: 0.5860 - val_loss: 0.6607 - val_accuracy: 0.5960

Epoch 2/5

45/45 [=====] - 1s 20ms/step - loss: 0.6494 - accuracy: 0.6167 - val_loss: 0.6488 - val_accuracy: 0.6180

Epoch 3/5

45/45 [=====] - 1s 14ms/step - loss: 0.6426 - accuracy: 0.6313 - val_loss: 0.6410 - val_accuracy: 0.6420

Epoch 4/5

45/45 [=====] - 1s 15ms/step - loss: 0.6303 - accuracy: 0.6473 - val_loss: 0.6332 - val_accuracy: 0.6340

Epoch 5/5

45/45 [=====] - 1s 16ms/step - loss: 0.6178 - accuracy: 0.6649 - val_loss: 0.6385 - val_accuracy: 0.6420

```
Out[32]: <keras.callbacks.History at 0x7ff973095940>
```

```
In [35]: y_pred = glove_dnn.predict(test_glove_features)
```

```
y_classes = np.argmax(y_pred, axis=1)
predictions = le.inverse_transform(y_classes)
# predictions = le.inverse_transform(y_pred)
```

63/63 [=====] - 0s 3ms/step

```
In [36]: meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_classes=['positive', 'negative'])
```

Model Performance metrics:

 Accuracy: 0.6475
 Precision: 0.6517
 Recall: 0.6475
 F1 Score: 0.6437

Model Classification report:

	precision	recall	f1-score	support
positive	0.67	0.54	0.60	981
negative	0.63	0.75	0.68	1019
accuracy			0.65	2000
macro avg	0.65	0.65	0.64	2000
weighted avg	0.65	0.65	0.64	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	534	447
negative	258	761